

L1. Simple Linear Regression (SLR)

- Overview of Statistical Regression Models and Machine Learning
- SLR : model set up
- SLR : least square estimate of β_0 and β_1
- SLR : mean and variance of β_0 and β_1

I. Overview of Statistical Regression Models

and supervised Machine Learning models.

Regression Models $\xrightarrow{\text{More data}}$ ML

\downarrow More features \downarrow

More about inferences

More about prediction

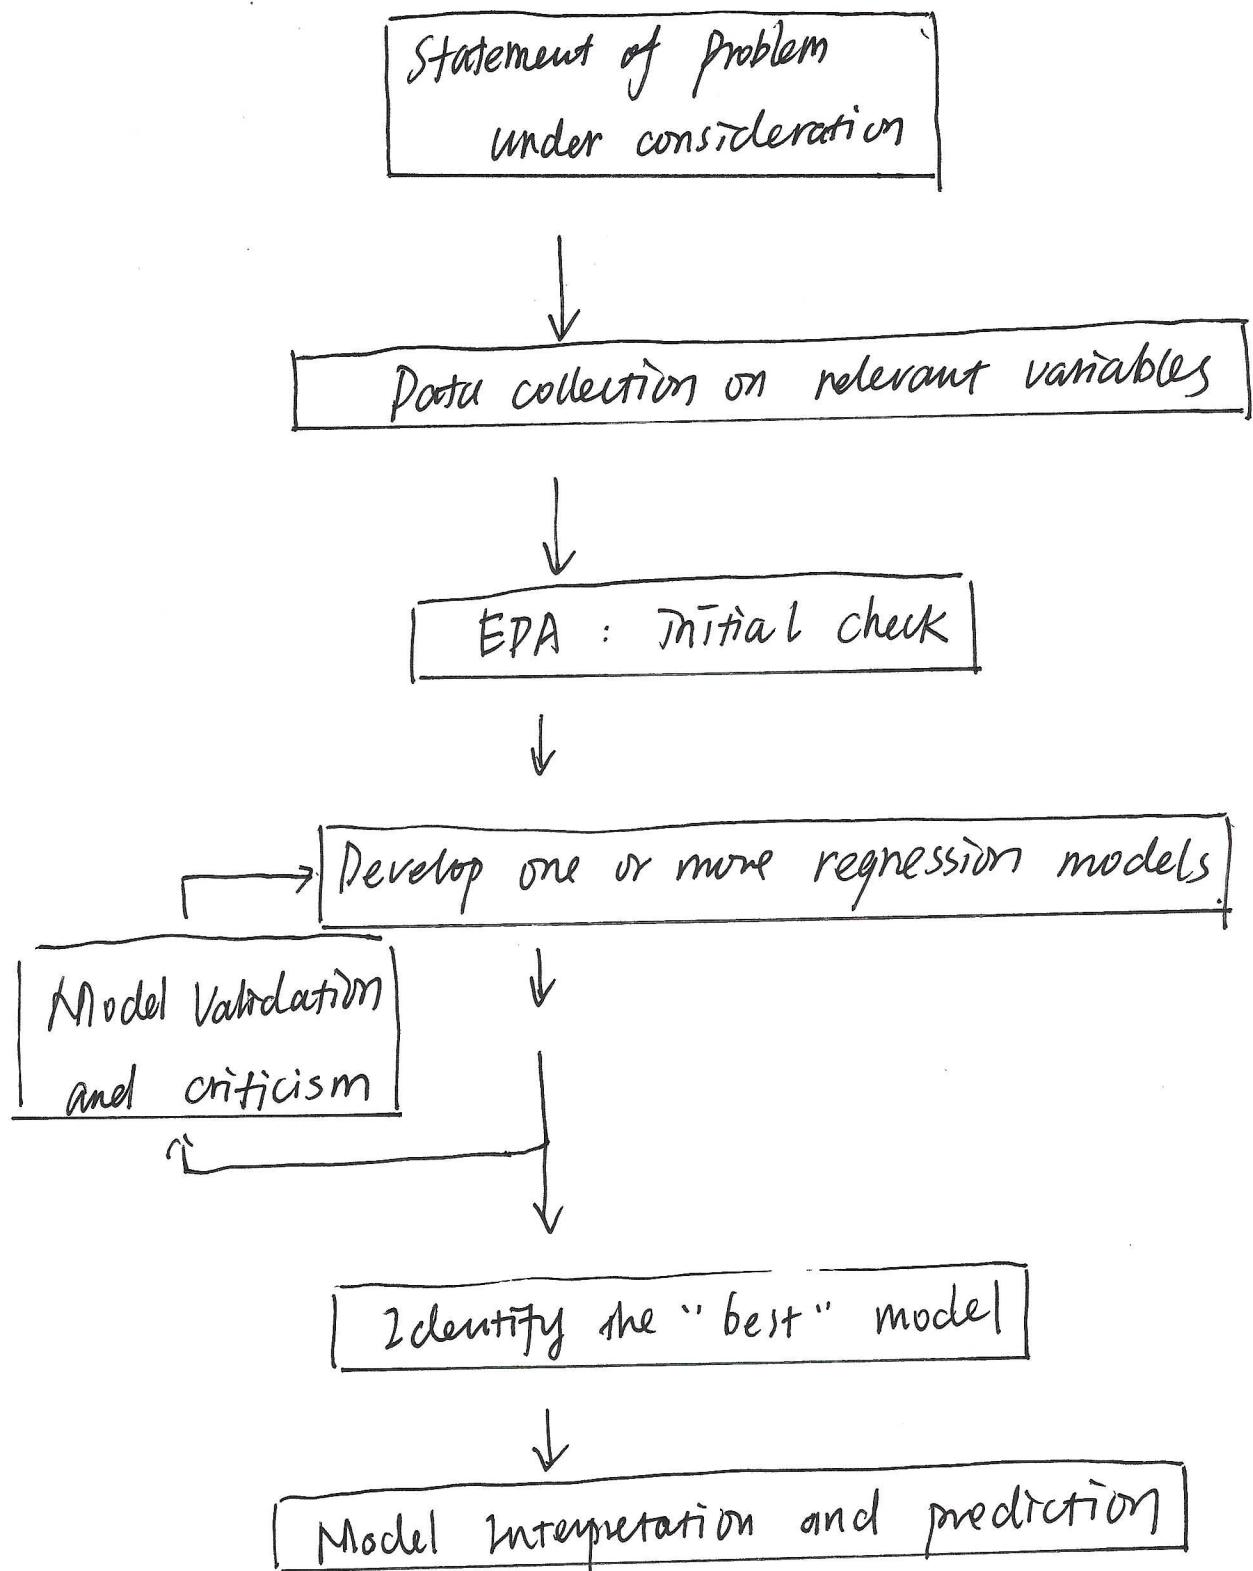
- relationship between variables
- significant feature
- inferences from the data.

- compare multiple approaches
- use train and test
- higher prediction accuracy



Foundation and Interpretation

Strategy of regression Models



II. SLR : Model set-up

1. Classic Simple Linear Regression Model

$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i, \quad i=1, \dots, n$$

- y_i : i^{th} value of the response variable y
- x_i : i^{th} value of the predictor / independent variable X
- n : total number of observations
- β_0 : intercept parameter (when $x=0$, $y=\beta_0$)
- β_1 : slope parameter (when $x \uparrow 1$, $y \uparrow \beta_1$)
- ε_i : random error

NOTE

1. Classic: related to model assumptions are "standard"
2. Simple: only one predictor X
3. Linear: y is linear to parameters β_0 and β_1

EXAMPLE: $y_i = \beta_0 + \beta_1^2 / \beta_0 x_i + \varepsilon_i$ — ^{non}linear

$y_i = \beta_0 + \beta_1 x_i^2 + \varepsilon_i$ — ~~non~~ linear.

2. Model Assumptions

- ε_i 's are the RANDOM error terms that satisfy

(i) $E(\varepsilon_i) = 0$

(ii) $\text{Var}(\varepsilon_i) = \sigma^2$ - constant

(iii) $\text{Cov}(\varepsilon_i, \varepsilon_j) = 0$ - no correlations between errors

error for one obs. can't predict the error
of another obs.

(iv) CLASSIC assumption :

$$\varepsilon_i \stackrel{\text{i.i.d.}}{\sim} N(0, \sigma^2)$$

- Not required for LSE

- needed for model Inference.

- X is a FIXED effect to Y

(i) X_i 's are treated as individual constants.

(ii) We don't need to worry about X 's distribution.

— There is a mixed effect model when you need to include random effects. Not discussed in this course.

3. Regression Function.

The regression function is the mean value $E(y_i)$

of the response variable y given x_i , the goal of
SLR is to estimate $E(y_i)$

Recall : $y_i = \beta_0 + \beta_1 x_i + \varepsilon_i$

$\uparrow \quad \underbrace{w}_{\text{fixed}} \quad \underbrace{w}_{\text{random}}$

- Regression Function:
$$\begin{aligned} E(y_i) &= E(\beta_0 + \beta_1 x_i) + E(\varepsilon_i) \\ &= (\beta_0 + \beta_1 x_i) + 0 \end{aligned}$$

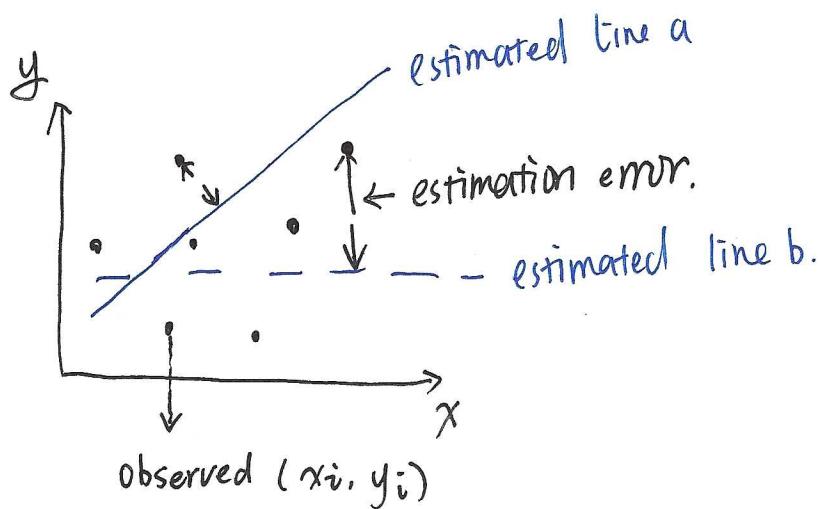
$$E(y_i) = \beta_0 + \beta_1 x_i$$

- $$\begin{aligned} \text{Var}(y_i) &= \text{Var}(\beta_0 + \beta_1 x_i + \varepsilon_i) = \text{Var}(\varepsilon_i) = \sigma^2 \end{aligned}$$
- $$\begin{aligned} \text{cov}(y_i, y_j) &= \text{cov}(\beta_0 + \beta_1 x_i + \varepsilon_i, \beta_0 + \beta_1 x_j + \varepsilon_j) \\ &= \text{cov}(\varepsilon_i, \varepsilon_j) \\ &= 0 \end{aligned}$$

III. SLR : least-square estimate of β_0 and β_1

Goal: we want to estimate $E(y_i) = \beta_0 + \beta_1 x_i \Leftrightarrow$ estimate β_0 and β_1

estimating principle: minimize the error between the observed y_i and estimated $\hat{y}_i = \hat{E}(y_i)$



1. Construct the deviation between the observed response y_i and the to-be-estimated linear mean component :

$$\varepsilon_i = y_i - (\beta_0 + \beta_1 x_i)$$

2. Construct a loss function ~~to~~ to evaluate the "overall" error; here the classic choice ~~is~~ is the sum of squares

$$S(\beta_0, \beta_1) = \sum_{i=1}^n \varepsilon_i^2 = \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2$$

3. Choose the values $(\hat{\beta}_0, \hat{\beta}_1)$ that minimize the loss function:

$$(\hat{\beta}_0, \hat{\beta}_1) = \operatorname{argmin} S(\beta_0, \beta_1)$$

Solution:

$$\begin{cases} \frac{\partial S}{\partial \beta_0} = 0 \\ \frac{\partial S}{\partial \beta_1} = 0 \end{cases} \Rightarrow \begin{cases} -2 \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i) = 0 \\ -2 \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i) x_i = 0 \end{cases}$$

$$\Rightarrow \begin{cases} \sum y_i - n \beta_0 - \beta_1 \sum x_i = 0 \\ \sum y_i x_i - \beta_0 \sum x_i - \beta_1 \sum x_i^2 = 0 \end{cases}$$

$$\Rightarrow \begin{cases} n \bar{y} - n \beta_0 - n \beta_1 \bar{x} = 0 \end{cases} \quad (1)$$

$$\begin{cases} \sum x_i y_i - n \beta_0 \bar{x} - \beta_1 \sum x_i^2 = 0 \end{cases} \quad (2)$$

Solve (1): $\beta_0 = \bar{y} - \beta_1 \bar{x}$

Plug into (2): $\sum x_i y_i - n \bar{x} (\bar{y} - \beta_1 \bar{x}) - \beta_1 \sum x_i^2 = 0$

$$\Rightarrow \beta_1 = \frac{\sum x_i y_i - n \bar{x} \bar{y}}{\sum x_i^2 - n \bar{x}^2} = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sum (x_i - \bar{x})^2}$$

Can you prove this is true?

Recall sample variance of x is $\sum_{i=1}^n (x_i - \bar{x})^2 / (n-1)$

sample covariance of x and y is $\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}) / (n-1)$

so some people denote β_1 as $\frac{s_{xy}}{s_{xx}}$

- The LSE of β_0 and β_1 are obtained as

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$$

$$\hat{\beta}_1 = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sum (x_i - \bar{x})^2} \quad \text{or} \quad \frac{s_{xy}}{s_{xx}}$$

Quick check for concavity:

$$\begin{cases} \frac{\partial^2 S}{\partial \beta_0^2} = 2n \\ \frac{\partial^2 S}{\partial \beta_1^2} = 2 \sum x_i^2 \\ \frac{\partial^2 S}{\partial \beta_0 \partial \beta_1} = 2 \sum x_i \end{cases} \Rightarrow H(\beta_0, \beta_1) = \begin{pmatrix} 2n & 2 \sum x_i \\ 2 \sum x_i & 2 \sum x_i^2 \end{pmatrix} \quad \text{Hessian Matrix}$$

$$\begin{aligned} \det(H) &= 4n \sum x_i^2 - 4 (\sum x_i)^2 \\ &= 4n \sum x_i^2 - 4(n\bar{x})^2 = 4n \sum x_i^2 - 4n^2 \bar{x}^2 \\ &= 4n (\sum x_i^2 - n\bar{x}) = 4n (\sum x_i^2 - 2n\bar{x} + n\bar{x}) \\ &= 4n (\sum x_i - 2\bar{x} \cdot \sum x_i + n\bar{x}) = 4n \sum (x_i - \bar{x})^2 > 0 \end{aligned}$$

Positive definite \rightarrow we have minimum.

IV. Mean of Variance of β_0 and β_1

How well the estimates performed can be evaluated by the mean and variance of the estimates:

- we want the estimate to be unbiased: $E(\hat{\beta}) = \beta$
- we want the estimate to be consistent: small variance.

For LSE of β_0 and β_1 , we have the following theory
to sum up the above:

Gauss-Markov Thm:

The least-square estimates $\hat{\beta}_0$ and $\hat{\beta}_1$ are unbiased,

$$\text{i.e. } E(\hat{\beta}_0) = \beta_0, \quad E(\hat{\beta}_1) = \beta_1$$

Moreover, $\hat{\beta}_0$ and $\hat{\beta}_1$ are BLUE

BLUE: Best Linear Unbiased Estimator.

- Best: smallest variance among all unbiased estimators of β_0 and β_1
- Linear: the estimator is a linear function of the data.

By proving the Gauss-Markov Thm, we can get the following:

$$1. E(\hat{\beta}_0) = \beta_0$$

$$2. E(\hat{\beta}_1) = \beta_1$$

$$3. \text{Var}(\hat{\beta}_0) = \sigma^2 \left(\frac{1}{n} + \frac{\bar{x}^2}{S_{xx}} \right)$$

$$4. \text{Var}(\hat{\beta}_1) = \frac{\sigma^2}{S_{xx}}$$

Proof. : we will start with $\hat{\beta}_1$

$$\text{recall } \hat{\beta}_1 = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sum (x_i - \bar{x})^2} = \frac{\sum (x_i - \bar{x})y_i}{\sum (x_i - \bar{x})^2}$$

Linear!!! $= \sum k_i y_i$

$$\text{where } k_i = \frac{x_i - \bar{x}}{\sum (x_i - \bar{x})^2}$$

Hint: $\sum (x_i - \bar{x})(y_i - \bar{y})$
 $= \sum x_i y_i - n \bar{x} \bar{y}$
 $= \sum (x_i - \bar{x}) y_i$

PAUSE

THREE FACTS ABOUT THE k_i 's

1) $\sum k_i = 0$

$$\sum k_i = \sum (x_i - \bar{x}) / \sum (x_i - \bar{x})^2 = \frac{\sum x_i - n \bar{x}}{\sum (x_i - \bar{x})^2} = \frac{n \bar{x} - n \bar{x}}{\sum (x_i - \bar{x})^2} = 0$$

2) $\sum k_i^2 = \frac{1}{\sum (x_i - \bar{x})^2}$

$$\sum k_i^2 = \sum \left(\frac{x_i - \bar{x}}{\sum (x_i - \bar{x})^2} \right)^2 = \frac{\sum (x_i - \bar{x})^2}{\left(\sum (x_i - \bar{x})^2 \right)^2} = \frac{1}{\sum (x_i - \bar{x})^2}$$

$$37 \quad \sum k_i x_i = 1$$

$$\sum k_i x_i = \sum (x_i - \bar{x}) x_i / \sum (x_i - \bar{x})^2 = \frac{\sum (x_i)^2 - n \bar{x}^2}{\sum (x_i - \bar{x})^2} \xrightarrow{\text{same}} = 1$$

pause ends

$$\text{Now: } \hat{\beta}_1 = \sum k_i y_i$$

$$\begin{aligned} E(\hat{\beta}_1) &= E(\sum k_i y_i) = \sum k_i E(y_i) = \sum k_i (\beta_0 + \beta_1 x_i) \\ &= \underbrace{\beta_0 \sum k_i}_{=0} + \underbrace{\beta_1 \sum k_i x_i}_{=1} \end{aligned}$$

UNBIASED!! $= \beta_1.$

$$\begin{aligned} \text{Var}(\hat{\beta}_1) &= \text{Var}(\sum k_i y_i) = \sum_{i=1}^n k_i^2 \text{Var}(y_i) \\ &= \sum k_i^2 \sigma^2 = \frac{\sigma^2}{\sum (x_i - \bar{x})^2} \quad \begin{matrix} \leftarrow \text{variance in } y \\ \leftarrow \text{variance in } x \end{matrix} \\ &\qquad \qquad \qquad S_{xx} \end{aligned}$$

Why Best? Will
Answer in MLR

$$E(\hat{\beta}_0) = E(\bar{y} - \bar{x} \hat{\beta}_1) = E(\bar{y}) - \bar{x} E(\hat{\beta}_1)$$

$$= \frac{1}{n} \sum E(y_i) - \bar{x} \beta_1$$

$$= \frac{1}{n} \sum (\beta_0 + \beta_1 x_i) - \bar{x} \beta_1 = \beta_0$$

UNBIASED!!

Now $\hat{\beta}_0$ is also linear, AKA. $\hat{\beta}_0 = \sum c_i y_i$

$$\hat{\beta}_0 = \bar{y} - \bar{x} \hat{\beta}_1 = \frac{1}{n} \sum y_i - \bar{x} \sum k_i y_i = \sum \left(\frac{1}{n} - \bar{x} k_i \right) y_i$$

$$\text{Var}(\hat{\beta}_0) = \left[\sum \left(\frac{1}{n} - \bar{x} k_i \right)^2 \right] \sigma^2$$

$$\text{Now, } \sum \left(\frac{1}{n} - \bar{x} k_i \right)^2$$

$$= \sum \left(\frac{1}{n^2} - \frac{2\bar{x} k_i}{n} + \bar{x}^2 k_i^2 \right)$$

$$= \cancel{\sum} \cancel{\left(\frac{1}{n^2} \right)}$$

$$= \frac{1}{n} - \frac{2\bar{x}}{n} \underbrace{\sum k_i}_{\substack{|| \\ 0}} + \bar{x}^2 \underbrace{\sum k_i^2}_{\substack{|| \\ \frac{1}{\sum (x_i - \bar{x})^2}}}$$

$$= \frac{1}{n} + \frac{\bar{x}^2}{\sum (x_i - \bar{x})^2}$$

$$\text{so } \text{Var}(\hat{\beta}_0) = \sigma^2 \left(\frac{1}{n} + \frac{\bar{x}^2}{S_{xx}} \right)$$

Summary and take-aways of this lecture

Classic Simple Linear Regression Model

$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i, \quad i=1, \dots, n$$

$$E(\varepsilon_i) = 0, \quad \text{Var}(\varepsilon_i) = \sigma^2, \quad \text{cov}(\varepsilon_i, \varepsilon_j) = 0$$

$$\varepsilon_i \stackrel{i.i.d.}{\sim} N(0, \sigma^2)$$

- $E(y_i) = \beta_0 + \beta_1 x_i$

$$\text{Var}(y_i) = \sigma^2$$

- $y_i |_{x=x_i} \sim N(\beta_0 + \beta_1 x_i, \sigma^2)$

Least-square Estimate of β_0 and β_1

$$\text{minimize } S(\beta_0, \beta_1) = \sum_{i=1}^n (y_i - (\beta_0 + \beta_1 x_i))^2$$

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$$

$$\hat{\beta}_1 = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sum (x_i - \bar{x})^2}$$

- The Estimate of $E(y_i) |_{x_i}$

is given by

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$$

- The point (\bar{x}, \bar{y}) lies on the fitted line:

$$\bar{y} = \hat{\beta}_0 + \hat{\beta}_1 \bar{x}$$

Gauss-Markov Thm:

The least-square estimators are the **BLUE** estimators of β_0 and β_1 with

$$E(\hat{\beta}_0) = \beta_0, \quad E(\hat{\beta}_1) = \beta_1$$

$$\text{Var}(\hat{\beta}_0) = \sigma^2 \left(\frac{1}{n} + \frac{\bar{x}^2}{\sum (x_i - \bar{x})^2} \right)$$

$$\text{Var}(\hat{\beta}_1) = \sigma^2 / \sum (x_i - \bar{x})^2$$

- How the Variance of the estimator are the smallest among linear unbiased estimators - MLR

- How to get σ^2 ?

Next lecture